

Deep Learning

Supervised

- * FFNN
- * CNN
- * RNN, LSTM
- * Forward preperation
- * Error / Loss Function
- * Gradient Descent
- * Back propagation
- * overfitting
- * Normalizing


Unsupervised

- * Autoencoders
- * GAN ✓

Reinforcement Learning ✓

- * Deep Q nets
- * Bellman's equation
- * Q-learning Algorithm
- * State, action learning

RL Trading Bot



Agent: Bot (RL Algorithm)

Actions: Buy, Sell, Hold

Environment: Stock market

State: # stocks you have
stock price
Account Balance

Reward: Profit / Loss

Value: maximize the profit

$$2000 + 2000 = 22000$$

example: current state $\rightarrow 10000, 2, 2000$ defined by function


\downarrow Buy 500 stock

Profit 105
(+ Reward)

new state $\rightarrow 10500, 2.01, 1000$ $2000 - (2 \times 500)$

$$21105 + 1000 = 22105$$

RL for Traffic Light control



Agent : RL Algorithm

Actions : Lt_1, Lt_2, Lt_3, Lt_4


Environment : Junction

State : Total waiting time
 wt_1, wt_2, wt_3, wt_4

Reward : depends on increasing/decreasing TWT

Value : minimize the TWT

House Example



Environment : House

Agent : RL Algorithm

States : current position
(Room #)


Action : movement

Reward : +/- Reward


Value : Reach #5

Reward matrix \rightarrow Define while exploring the environment

		Actions					
		0	1	2	3	4	5
States	0	-1	-1	-1	-1	0	-1
	1	-1	-1	-1	0	-1	100
	2	-1	-1	-1	0	-1	-1
	3	-1	0	0	-1	0	-1
	4	0	-1	-1	-1	-1	100
	5	-1	-1	-1	-1	0	100



RL Trading Bot



Agent: Bot (RL Algorithm)

Actions: Buy, Sell, Hold

Environment: Stock market

State: # stocks you have

stock price

Account Balance

Reward: Profit/Loss

Value: maximize the profit

Reward matrix

state space

States	$S_1 (1000, 2, 500)$
	$S_2 (1000, 3, 200)$
	$S_3 (1500, 5, 100)$
	\vdots
	S_n

Actions

Buy Sell Hold ← action space

0	100	70
-1	0	100
0	100	20
\vdots		

Q matrix

$\alpha = 1$

0

$$Q = R + \gamma \cdot \max_a Q$$

$$Q(\text{state}, \text{action}) \leftarrow (1 - \alpha)Q(\text{state}, \text{action}) + \alpha \left(\text{reward} + \gamma \max_a Q(\text{next state}, \text{all actions}) \right)$$

Initialize Q

Choose action from Q

Perform action

Measure Reward

Update Q

R matrix

	Action					
State	0	1	2	3	4	5
0	-1	-1	-1	-1	0	-1
1	-1	-1	-1	0	-1	100
2	-1	-1	-1	0	-1	-1
3	-1	0	0	-1	0	-1
4	0	-1	-1	0	-1	100
5	-1	0	-1	-1	0	100

The -1's in the table represent null values

Q matrix

Actions


	0	1	2	3	4	5
0	-1	-1	-1	-1	0	-1
1	-1	-1	-1	0	-1	100
2	-1	-1	-1	0	-1	-1
3	-1	0	0	-1	0	-1
4	0	-1	-1	-1	-1	100
5	-1	-1	-1	-1	0	100



$$0 \leq \gamma \leq 1$$

$$\gamma = 0.8$$

$$Q(s, A) \leftarrow \text{reward} + \gamma \max_a Q(\text{next state}, \text{all actions})$$



R matrix

	Action					
State	0	1	2	3	4	5
0	-1	-1	-1	-1	0	-1
1	-1	-1	-1	0	-1	100
2	-1	-1	-1	0	-1	-1
3	-1	0	0	-1	0	-1
4	0	-1	-1	0	-1	100
5	-1	0	-1	-1	0	100

The -1's in the table represent null values

Q matrix

	Actions					
	0	1	2	3	4	5
0	0	0	0	0	80	0
1	0	0	0	0	0	0
2	0	0	0	0	0	0
3	0	0	0	0	80	0
4	0	0	0	0	0	100
5	0	0	0	0	0	0

$$Q(2,3) = 0 + 0.8(0) = 0$$


$$Q(3,4) = 0 + 0.8(0) = 0$$

$$Q(4,5) = 100 + 0.8(0) = 100$$

$$Q(0,4) = 0 + 0.8(100) = 80$$

$$Q(4,3) = 0 + 0.8(0) = 0$$


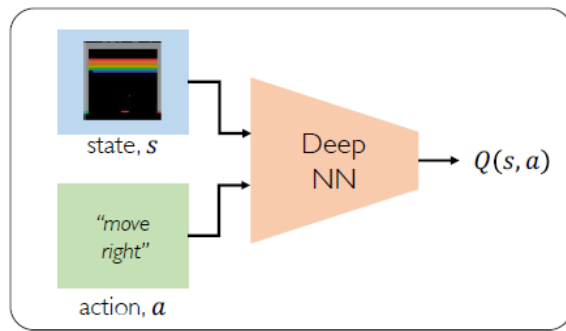
$$Q(3,4) = 0 + 0.8(100) = 80$$



trained Q

	Actions					
	0	1	2	3	4	5
0	0.	0.	0.	0.	80.	0.
1	0.	0.	0.	64.	0.	100.
2	0.	0.	0.	64.	0.	0.
3	0.	80.	51.2	0.	80.	0.
4	64.	0.	0.	64.	0.	100.
5	0.	80.	0.	0.	80.	100.

How can we use deep neural networks to model Q-functions?



$$\mathcal{L} = \mathbb{E} \left[\left\| \overbrace{\left(r + \gamma \max_{a'} Q(s', a') \right)}^{\text{target}} - \overbrace{Q(s, a)}^{\text{predicted}} \right\|^2 \right]$$

↑
actual